

Lesson 10 Demo 03

Creating an ECS Cluster with Auto Scaling

Objective: To demonstrate how to create an ECS cluster with auto scaling and Container Insights, and how to run a task definition for an Nginx container on the cluster

Tools required: AWS Management Console

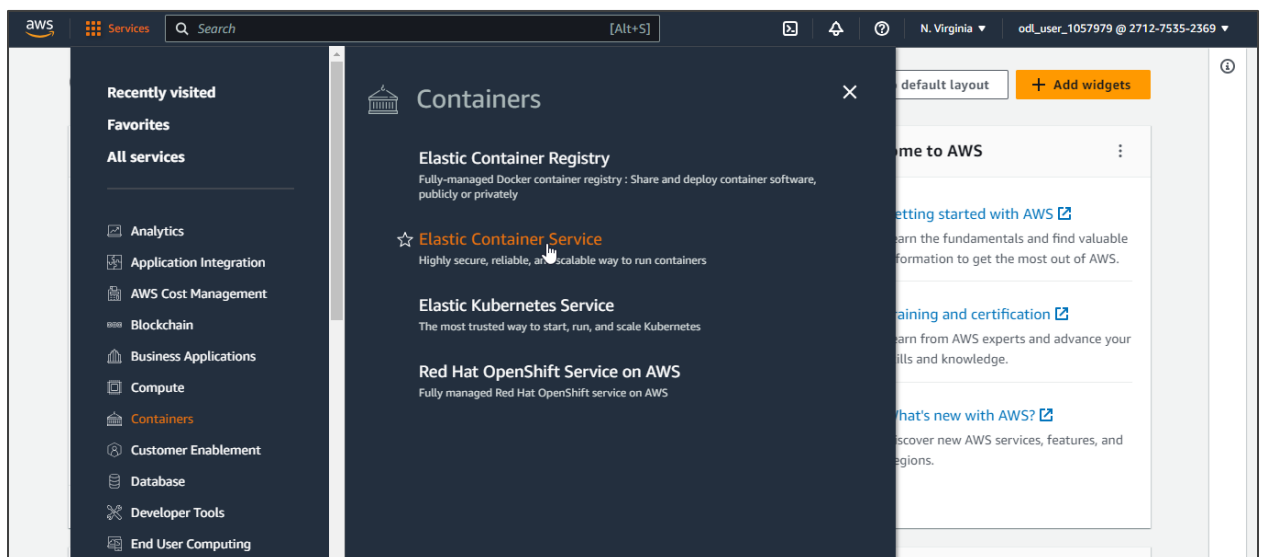
Prerequisites: A key pair for an EC2 instance must be created.

Steps to be followed:

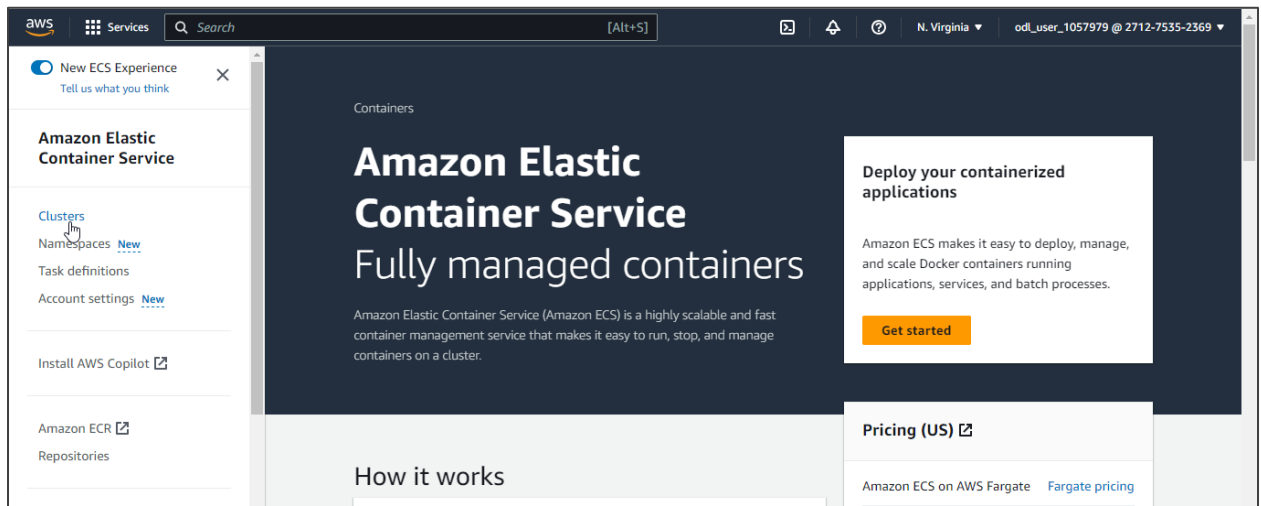
1. Create an ECS cluster
2. Create a task definition
3. Run the task definition on the cluster

Step 1: Create an ECS cluster

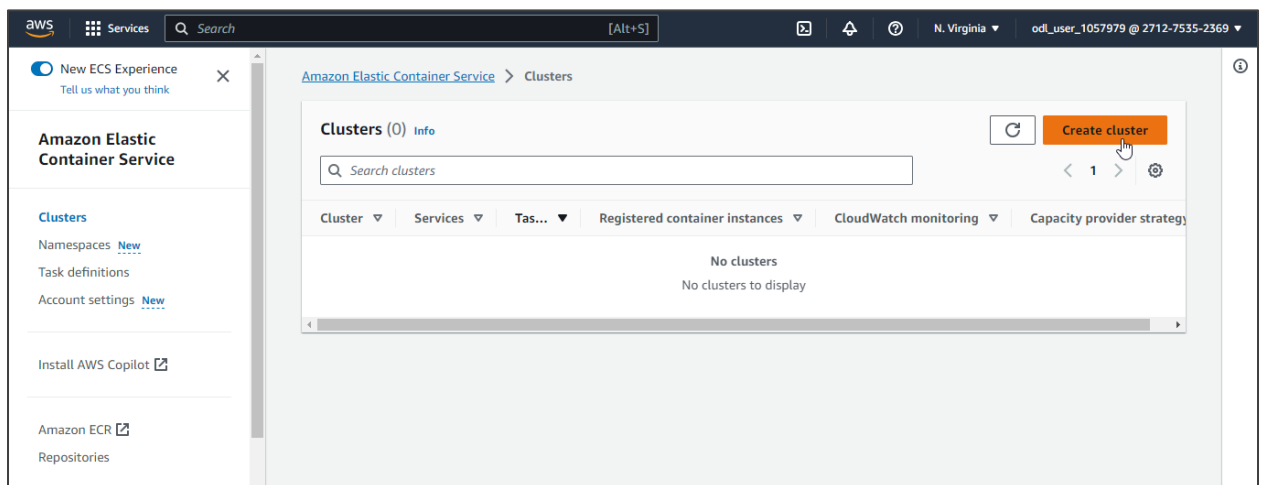
1.1 In the AWS Console, select **Services** and then choose **ECS** under **Containers**



1.2 In the navigation pane, choose **Clusters**



1.3 Choose **Create cluster**



1.4 Create an ECS cluster with the settings shown in the screenshots:

Create cluster [Info](#)

An Amazon ECS cluster groups together tasks, and services, and allows for shared capacity and common configurations. All of your tasks, services, and capacity must belong to a cluster.

Cluster configuration

Cluster name

There can be a maximum of 255 characters. The valid characters are letters (uppercase and lowercase), numbers, hyphens, and underscores.

Default namespace - *optional*
 Select the namespace to specify a group of services that make up your application. You can overwrite this value at the service level.

▼ Infrastructure [Info](#)

Your cluster is automatically configured for AWS Fargate (serverless) with two capacity providers. Add Amazon EC2 instances, or external instances using ECS Anywhere.

[Customized](#)

☐ AWS Fargate (serverless)
 Pay as you go. Use if you have tiny, batch, or burst workloads or for zero maintenance overhead. The cluster has Fargate and Fargate Spot capacity providers by default.

☒ Amazon EC2 instances
 Manual configurations. Use for large workloads with consistent resource demands.

Auto Scaling group (ASG) [Info](#)
 Use Auto Scaling groups to scale the Amazon EC2 instances in the cluster.

Operating system/Architecture
 Choose the Windows operating system or Linux architecture for your instance.

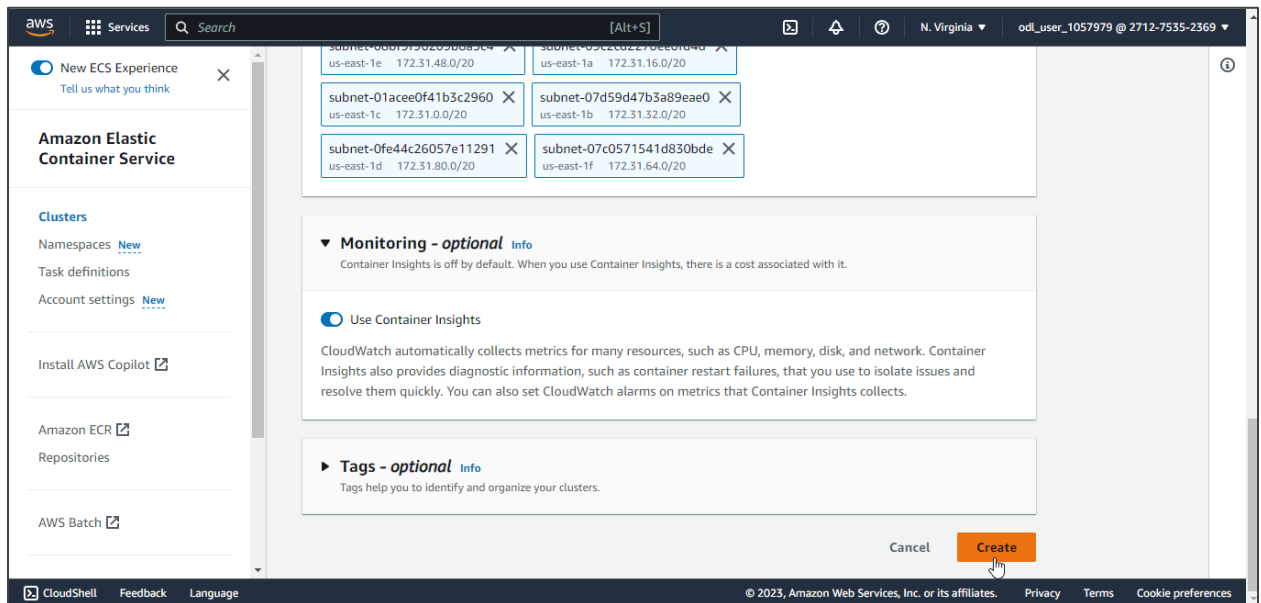
EC2 instance type
 Choose based on the workloads you plan to run on this cluster.
 [Free tier eligible](#)
 i386, x86_64
 1 vCPU 1 GiB Memory

Desired capacity
 Specify the number of instances to launch in your cluster.
 Minimum: Maximum:

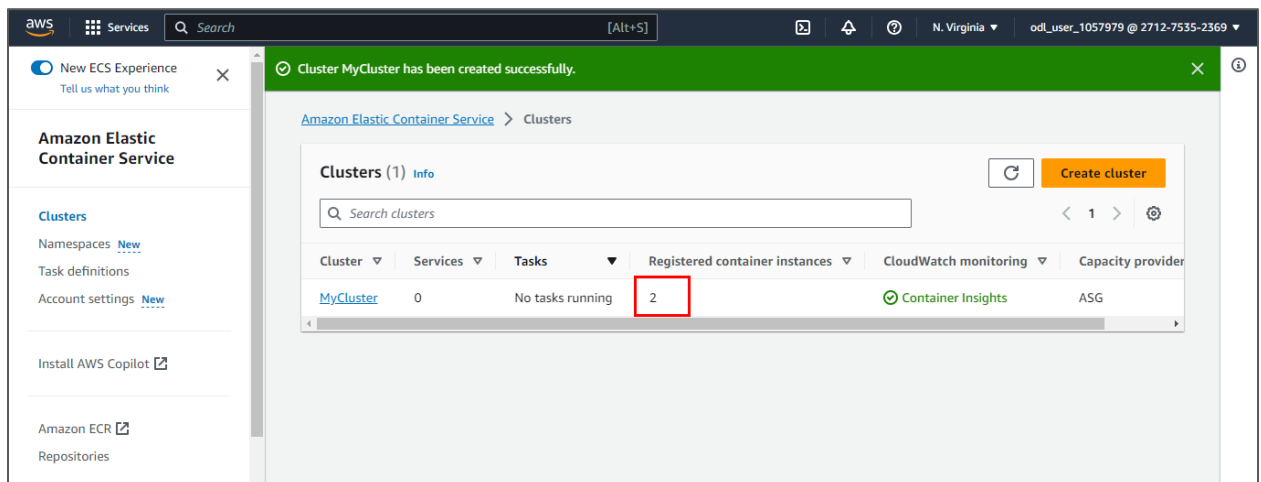
SSH Key pair
 Create a key pair in the EC2 console, consisting of a private key and a public key, that you use to prove your identity when connecting to an instance.

☐ External instances using ECS Anywhere
 Manual configurations. Use to add data center compute.

1.5 Click Create

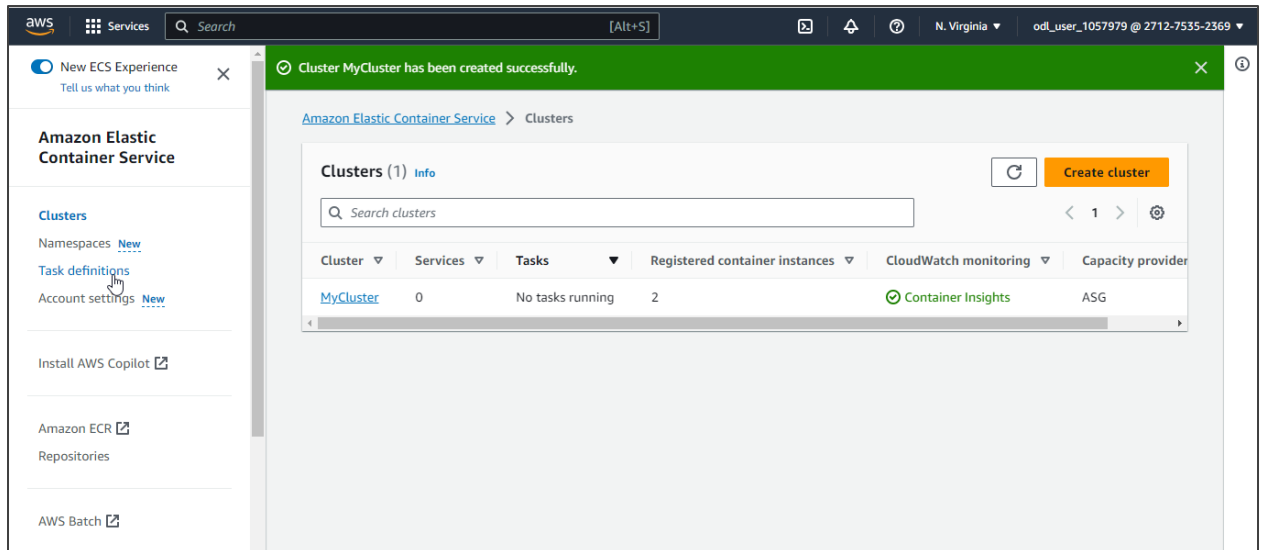


Note: Please wait for a few minutes until the EC2 instances are registered on the cluster.

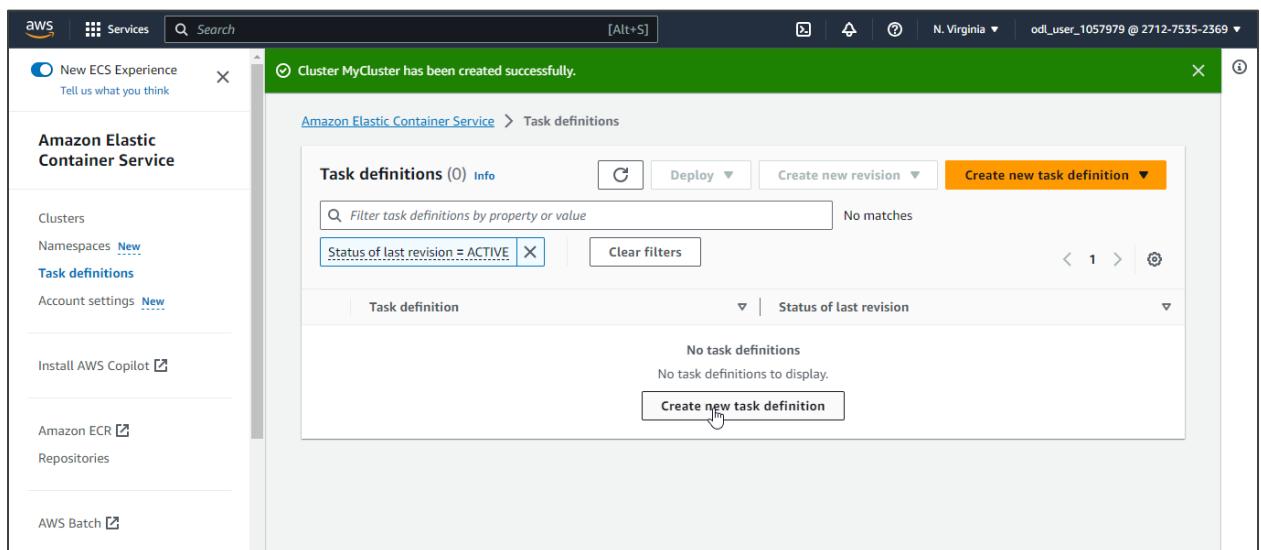


Step 2: Create a task definition

2.1 In the navigation pane, choose Task definitions



2.2 Select Create new task definition



2.3 Create the task definition with the settings shown in the screenshots:

Create new task definition [Info](#)

Task definition configuration

Task definition family [Info](#)
Specify a unique task definition family name.

Up to 255 letters (uppercase and lowercase), numbers, hyphens, and underscores are allowed.

▼ Infrastructure requirements
Specify the infrastructure requirements for the task definition.

Launch type [Info](#)
Selection of the launch type will change task definition parameters.

☐ AWS Fargate
Serverless compute for containers.

☒ Amazon EC2 instances
Self-managed infrastructure using Amazon EC2 instances.

Container - 1 [Info](#) Essential container Remove

Container details
Specify a name, container image, and whether the container should be marked as essential. Each task definition must have at least one essential container.

Name Image URI Essential container

Private registry [Info](#)
Store credentials in Secrets Manager, and then use the credentials to reference images in private registries.
☐ Private registry authentication

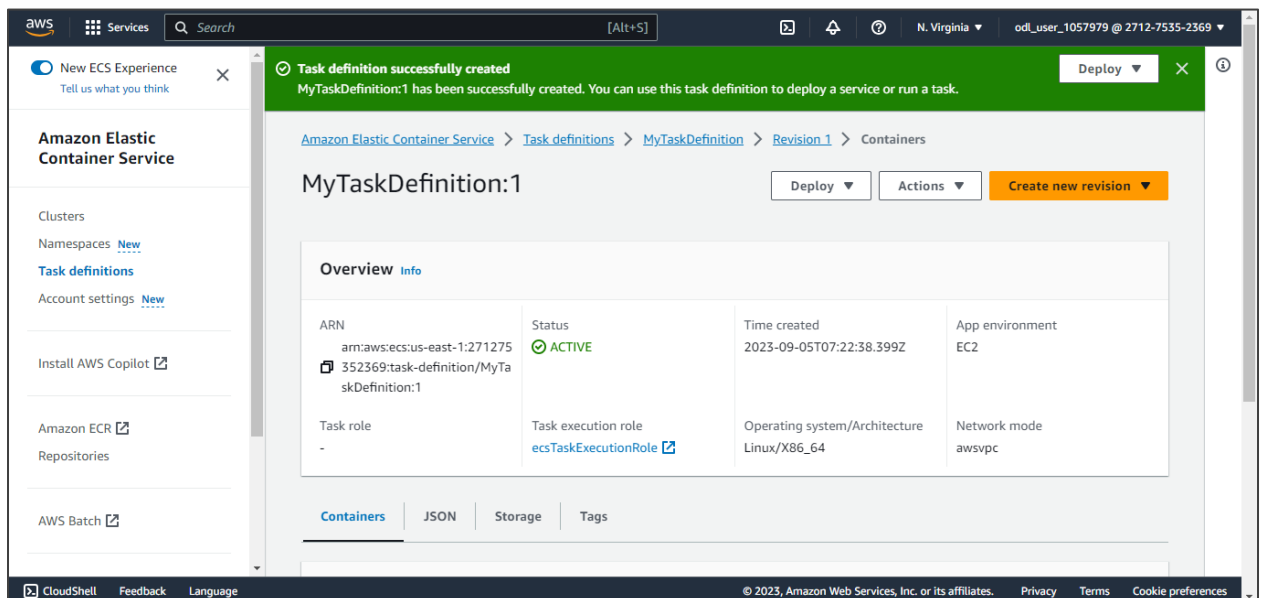
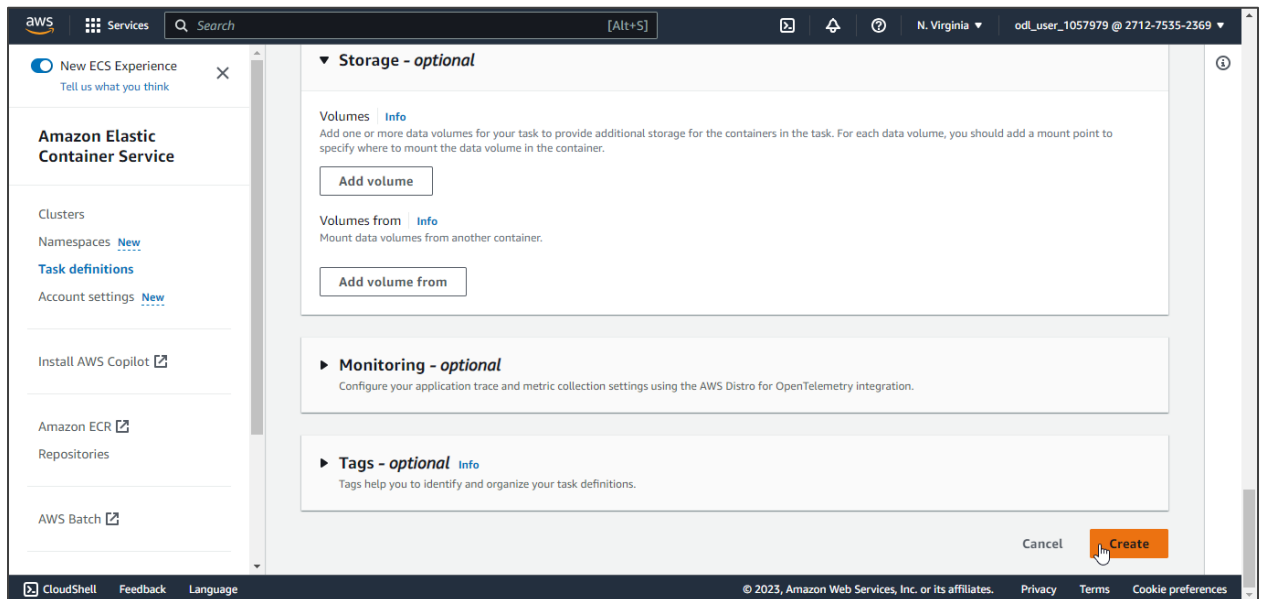
Port mappings [Info](#)
Add port mappings to allow the container to access ports on the host to send or receive traffic. Any changes to port mappings configuration impacts the associated service connect settings.

Container port	Protocol	Port name	App protocol	
<input type="text" value="80"/>	<input type="text" value="TCP"/>	<input type="text" value="container-80-tcp"/>	<input type="text" value="HTTP"/>	Remove

Add more port mappings

Read only root file system [Info](#)
When this parameter is turned on, the container is given read-only access to its root file system.
☐ Read only

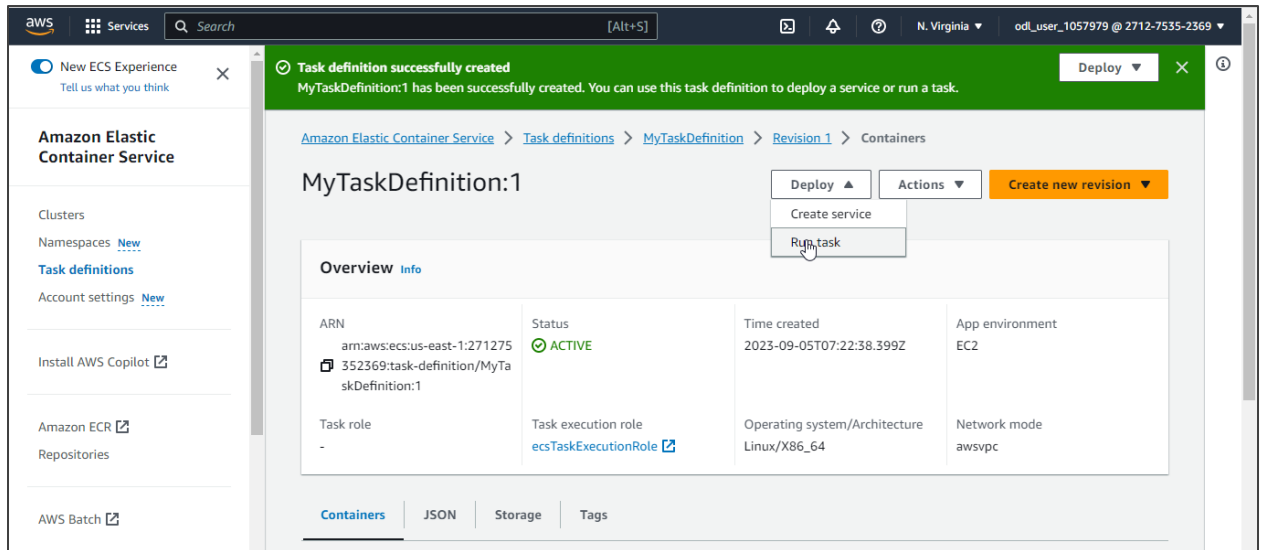
2.4 Click Create



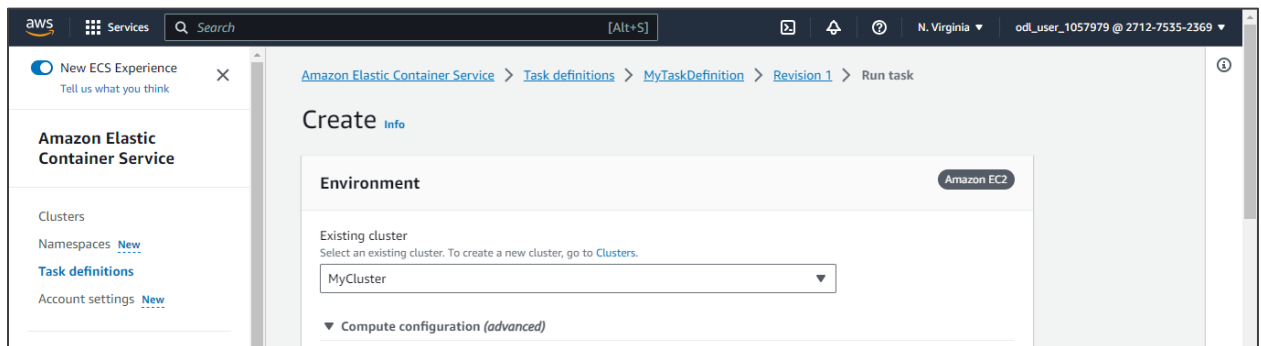
The task definition is created successfully.

Step 3: Run the task definition on the cluster

3.1 To run the task definition, select **Deploy** and click **Run task**



3.2 Choose the cluster that you created in the previous step



3.3 Configure the other settings as shown in the screenshots:

The screenshot shows the AWS Management Console for the Amazon Elastic Container Service. The left sidebar contains the 'Amazon Elastic Container Service' section with links to Clusters, Namespaces, Task definitions, Account settings, Install AWS Copilot, Amazon ECR, and AWS Batch. The main content area is divided into two sections: 'Compute options' and 'Deployment configuration'.

Compute options

To ensure task distribution across your compute types, use appropriate compute options.

☐ Capacity provider strategy
Specify a launch strategy to distribute your tasks across one or more capacity providers.

☒ Launch type
Launch tasks directly without the use of a capacity provider strategy.

Launch type

Select either managed capacity (Fargate), or custom capacity (EC2 or user-managed, External instances). External instances are registered to your cluster using the ECS Anywhere capability.

EC2

Deployment configuration

Application type

Specify what type of application you want to run.

☐ Service
Launch a group of tasks handling a long-running computing work that can be stopped and restarted. For example, a web application.

☒ Task
Launch a standalone task that runs and terminates. For example, a batch job.

The screenshot shows the AWS Management Console for the Amazon Elastic Container Service, specifically the 'Task definition' section. The left sidebar is the same as the previous screenshot. The main content area is titled 'Task definition' and includes the following sections:

Task definition

Select an existing task definition. To create a new task definition, go to Task definitions.

☐ Specify the revision manually
Manually input the revision instead of choosing from the 100 most recent revisions for the selected task definition family.

Family

MyTaskDefinition

Revision

1 (LATEST)

Desired tasks

Specify the number of tasks to launch.

1

Task group

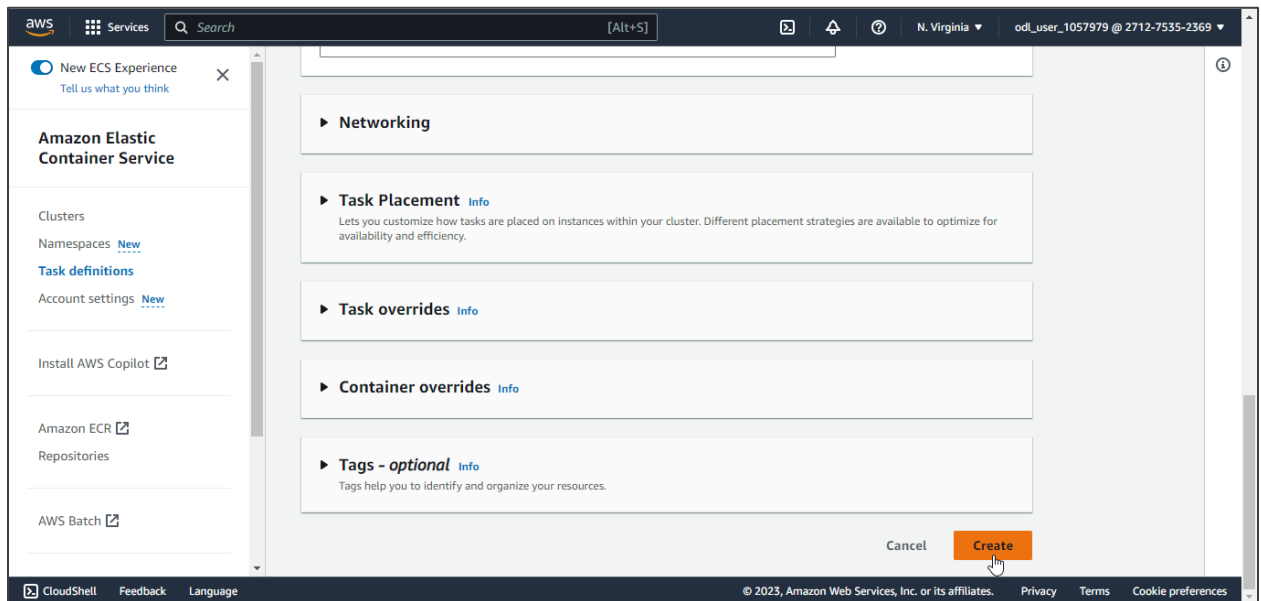
All tasks with the same task group name are considered as a set when performing spread placement.

Networking

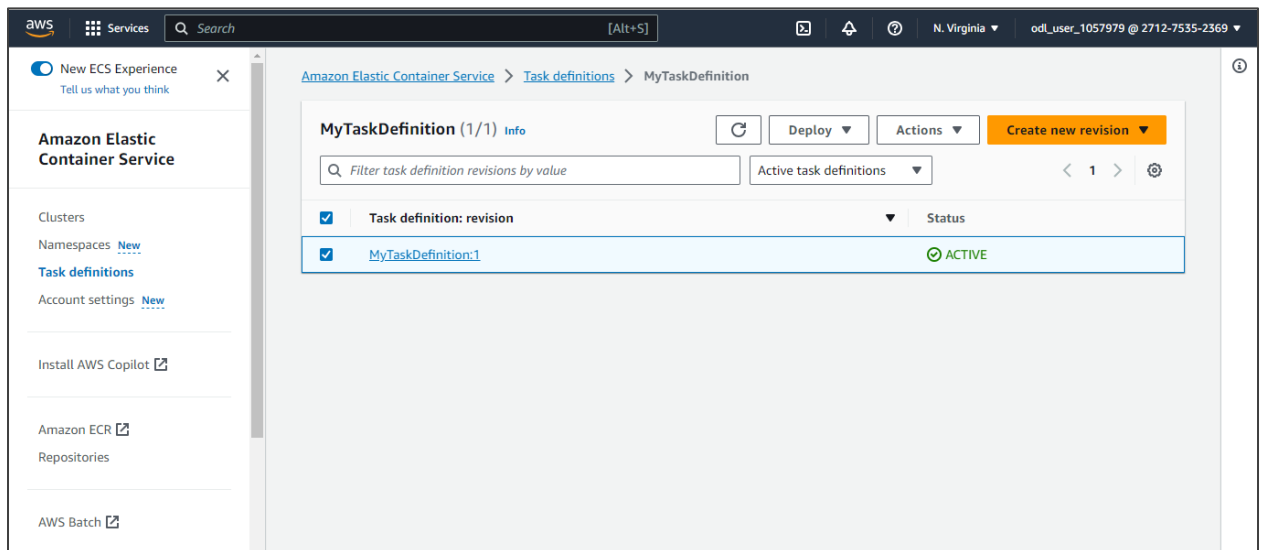
Task Placement

Lets you customize how tasks are placed on instances within your cluster. Different placement strategies are available to optimize for availability and efficiency.

3.4 Choose Create



The task definition is successfully executed.



The screenshot displays the AWS Management Console for an Amazon ECS cluster named 'MyCluster'. The cluster is in an 'Active' state. The 'Tasks' tab is selected, showing a single task with ID '538ff...' in a 'Running' state. The task is associated with the 'MyCluster' and has a health status of 'Unknown'.

Task	Last status	Desired st...	T...	Rev...	Health sta...	Started at	Container instan...
538ff...	Provisioning	Running	MyI...	1	Unknown	-	-

As you can see in the screenshot above, a task has been added to **MyCluster** and triggered successfully.

By following these steps, you have successfully created an ECS cluster with auto scaling and enabled Container Insights. You have also created a task definition for a Nginx container and run it on the cluster.