

Lesson 10 Demo 05

Running Task on a Fargate Cluster

Objective: To run a task on a Fargate cluster

Tools required: An AWS account

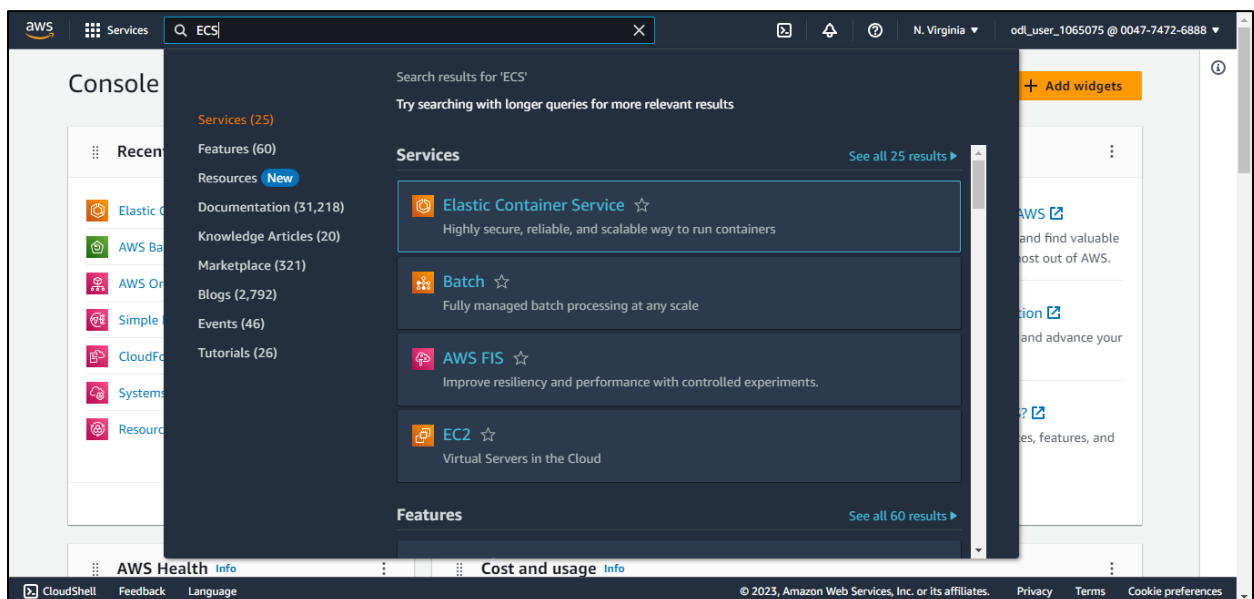
Prerequisites: None

Steps to be followed:

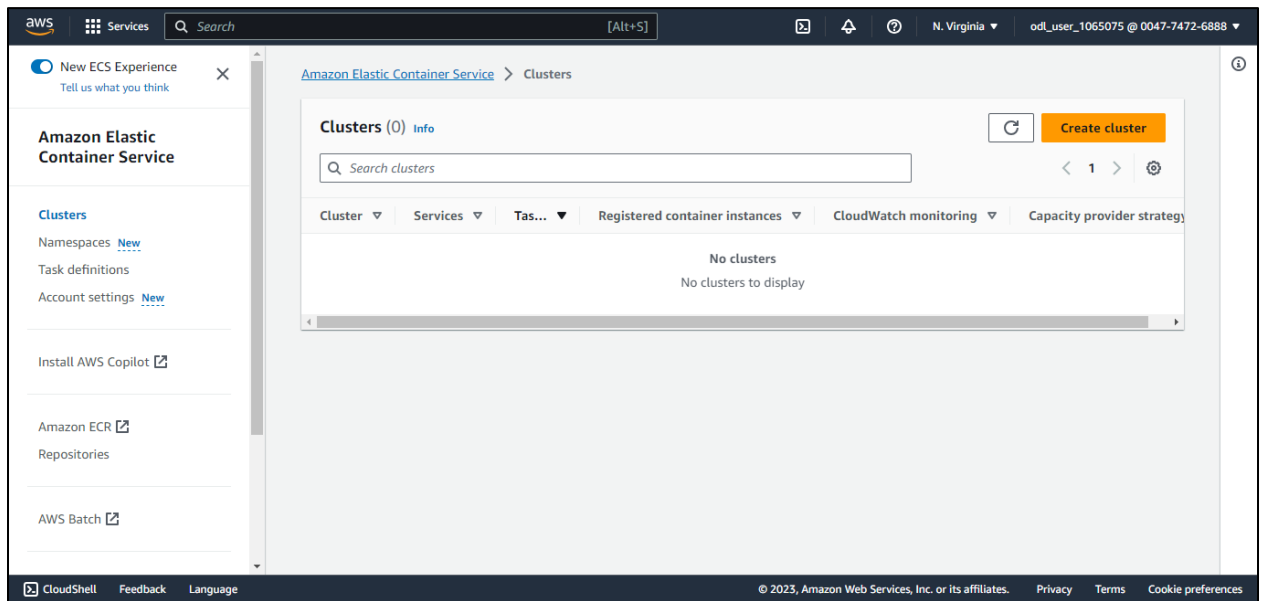
1. Create a Fargate cluster
2. Create a task definition
3. Run the Fargate Cluster

Step 1: Create a Fargate cluster

1.1 In the **AWS Management Console**, search for **ECS** and select **Elastic Container Service**



1.2 On the left panel of the ECS console, click on **Clusters** and **Create cluster**



1.3 In the **Cluster configuration**, perform the following:

- Enter an arbitrary name for the cluster under the **Cluster name**
- In the infrastructure, specify **AWS Fargate (serverless)**
- Let the other settings stay at default settings and click on **Create**

Create cluster [Info](#)

An Amazon ECS cluster groups together tasks, and services, and allows for shared capacity and common configurations. All of your tasks, services, and capacity must belong to a cluster.

Cluster configuration

Cluster name
simplifargate

There can be a maximum of 255 characters. The valid characters are letters (uppercase and lowercase), numbers, hyphens, and underscores.

Default namespace - *optional*
Select the namespace to specify a group of services that make up your application. You can overwrite this value at the service level.
simplifargate

▼ Infrastructure [Info](#)
Your cluster is automatically configured for AWS Fargate (serverless) with two capacity providers. Add Amazon EC2 instances, or external instances using ECS Anywhere.
Serverless

▼ Infrastructure [Info](#)
Your cluster is automatically configured for AWS Fargate (serverless) with two capacity providers. Add Amazon EC2 instances, or external instances using ECS Anywhere.
Serverless

☒ **AWS Fargate (serverless)**
Pay as you go. Use if you have tiny, batch, or burst workloads or for zero maintenance overhead. The cluster has Fargate and Fargate Spot capacity providers by default.

☐ **Amazon EC2 instances**
Manual configurations. Use for large workloads with consistent resource demands.

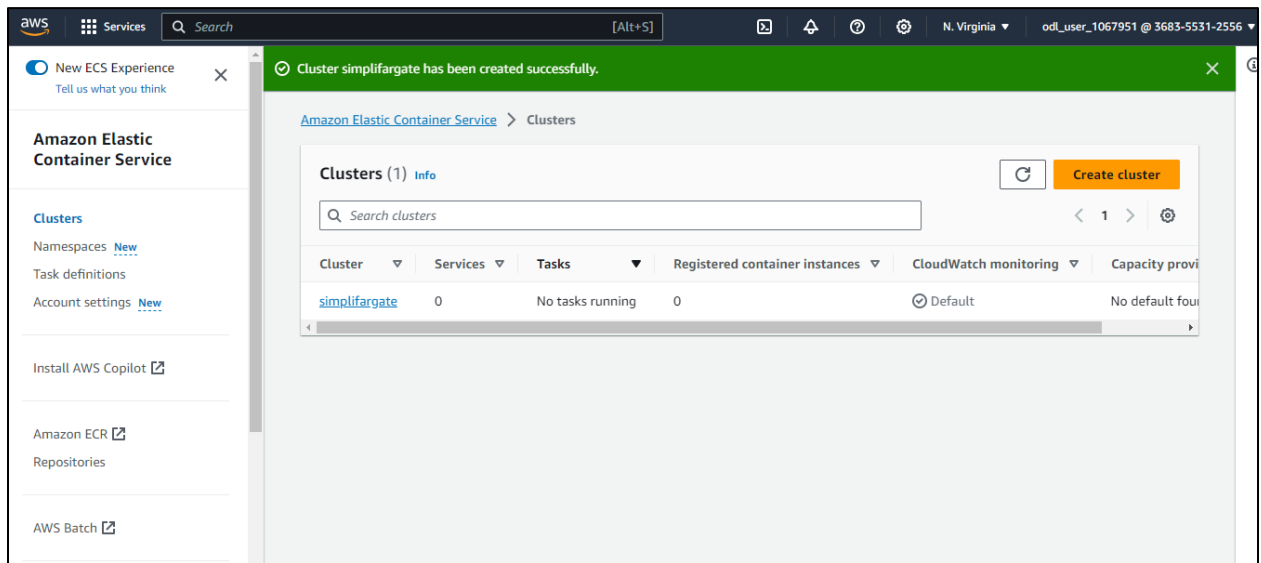
☐ **External instances using ECS Anywhere**
Manual configurations. Use to add data center compute.

► Monitoring - optional [Info](#)
Container Insights is off by default. When you use Container Insights, there is a cost associated with it.

► Tags - optional [Info](#)
Tags help you to identify and organize your clusters.

Cancel **Create**

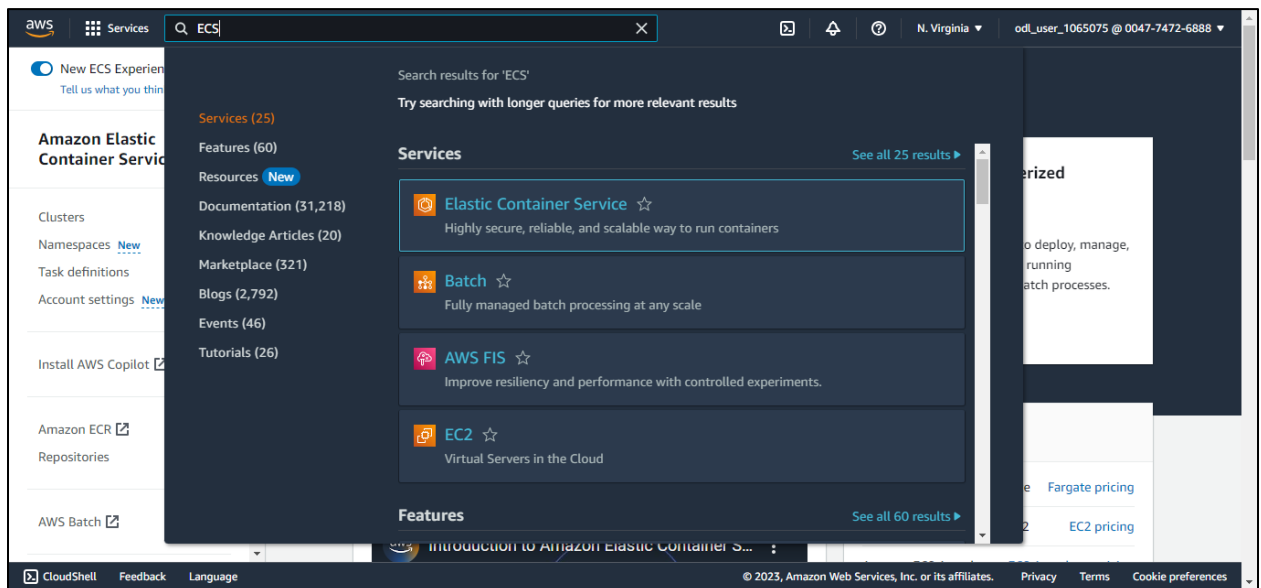
1.4 Verify the cluster creation as shown below:



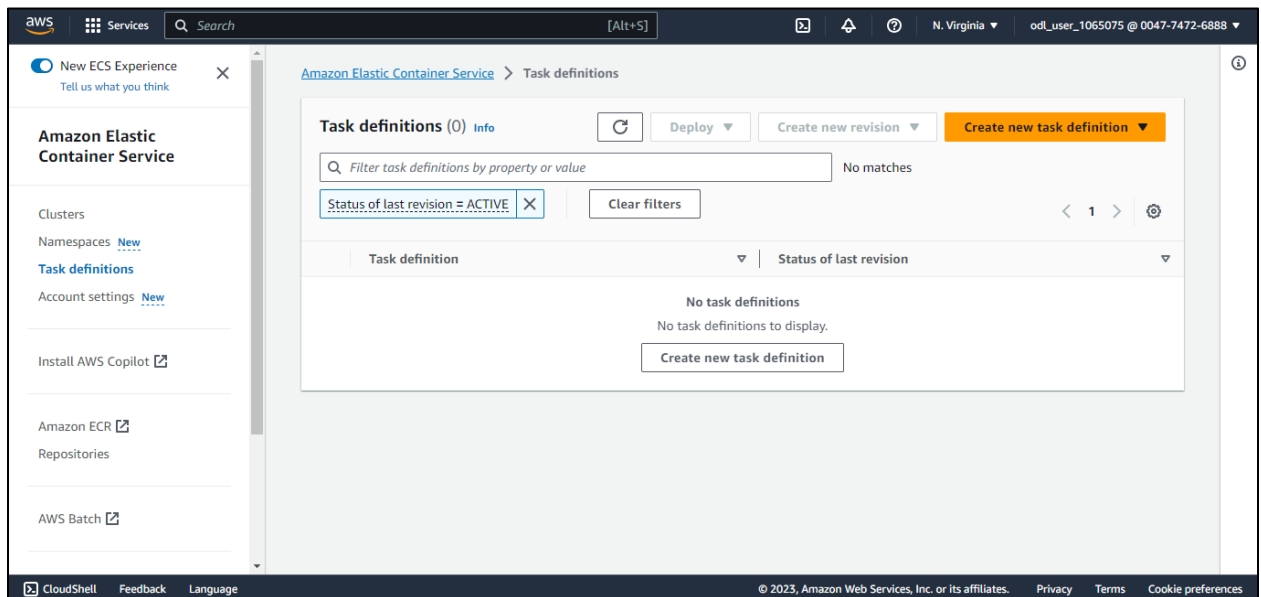
Note: Do not close the above tab. It will be necessary for reference. ECS Cluster will be created.

Step 2: Create a task definition

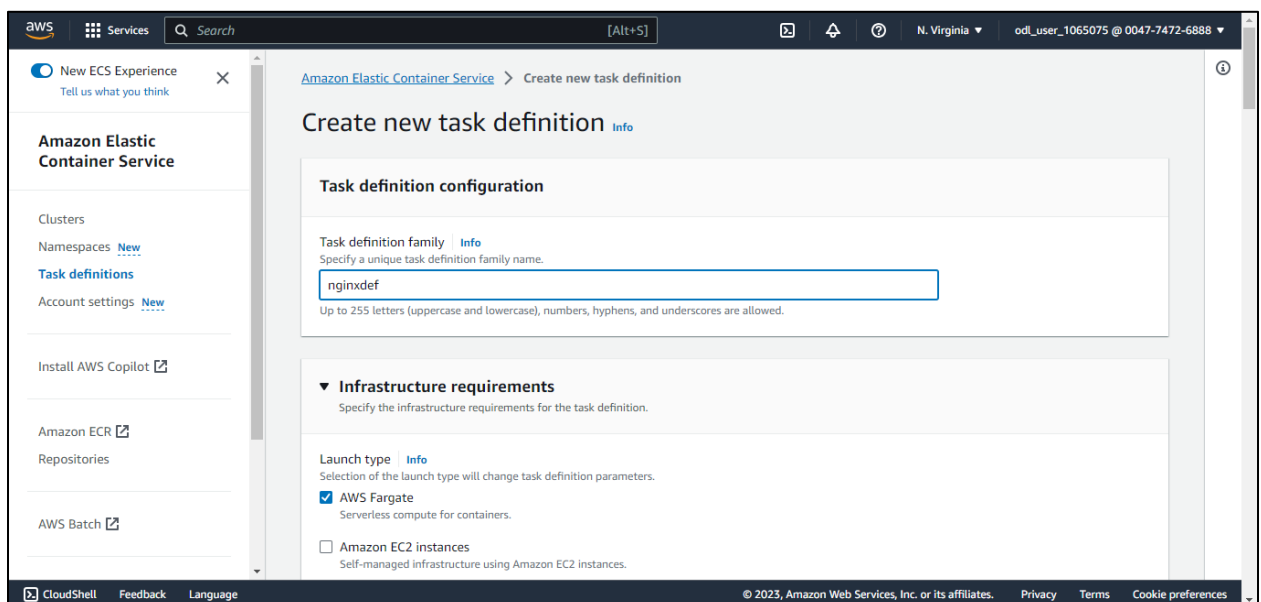
2.1 In the AWS Management Console, search for ECS and select Elastic Container Service



2.2 On the left panel of the ECS console, click on **Task definitions** and on **Create new task definition**



2.3 In the **Task definition configuration** page, specify task definition family = **nginxdef**, **Launch type** = **AWS Fargate**, **CPU** = **0.25 vCPU**, and **Memory** = **0.5 GB**



Infrastructure requirements
Specify the infrastructure requirements for the task definition.

Launch type [Info](#)
Selection of the launch type will change task definition parameters.

☒ **AWS Fargate**
Serverless compute for containers.

☐ **Amazon EC2 instances**
Self-managed infrastructure using Amazon EC2 instances.

OS, Architecture, Network mode
Network mode is used for tasks and is dependent on the compute type selected.

Operating system/Architecture [Info](#)
Linux/X86_64

Network mode [Info](#)
awsvpc

Task size [Info](#)
Specify the amount of CPU and memory to reserve for your task.

CPU .25 vCPU **Memory** .5 GB

Task roles - conditional

2.4 For Container-1 details, enter **Name = nginx** and **Image URI = public.ecr.aws/nginx/nginx:1.25**

Container - 1 [Info](#) Essential container Remove

Container details
Specify a name, container image, and whether the container should be marked as essential. Each task definition must have at least one essential container.

Name nginx **Image URI** public.ecr.aws/nginx/nginx:1.25 **Essential container** Yes

Private registry [Info](#)
Store credentials in Secrets Manager, and then use the credentials to reference images in private registries.

☐ **Private registry authentication**

Port mappings [Info](#)
Add port mappings to allow the container to access ports on the host to send or receive traffic. Any changes to port mappings configuration impacts the associated service connect settings.

| Container port | Protocol | Port name | App protocol | |
|----------------|----------|--------------|--------------|---------------------|
| 80 | TCP | nginx-80-tcp | HTTP | Remove |

Add more port mappings

Read only root file system [Info](#)
When this parameter is turned on, the container is given read-only access to its root file system.

☐ **Read only**

Resource allocation limits - conditional [Info](#)

2.5 Leave other options default and click **Create**

The screenshot shows the 'Create new task definition' page in the AWS Management Console. The left sidebar contains the 'Amazon Elastic Container Service' menu with options like Clusters, Namespaces, Task definitions, and Account settings. The main content area has a 'Volumes' section with an 'Add volume' button, a 'Volumes from' section with an 'Add volume from' button, and two optional sections: 'Monitoring - optional' and 'Tags - optional'. At the bottom right, there are 'Cancel' and 'Create' buttons. The 'Create' button is highlighted in orange.

The task definition has been created successfully:

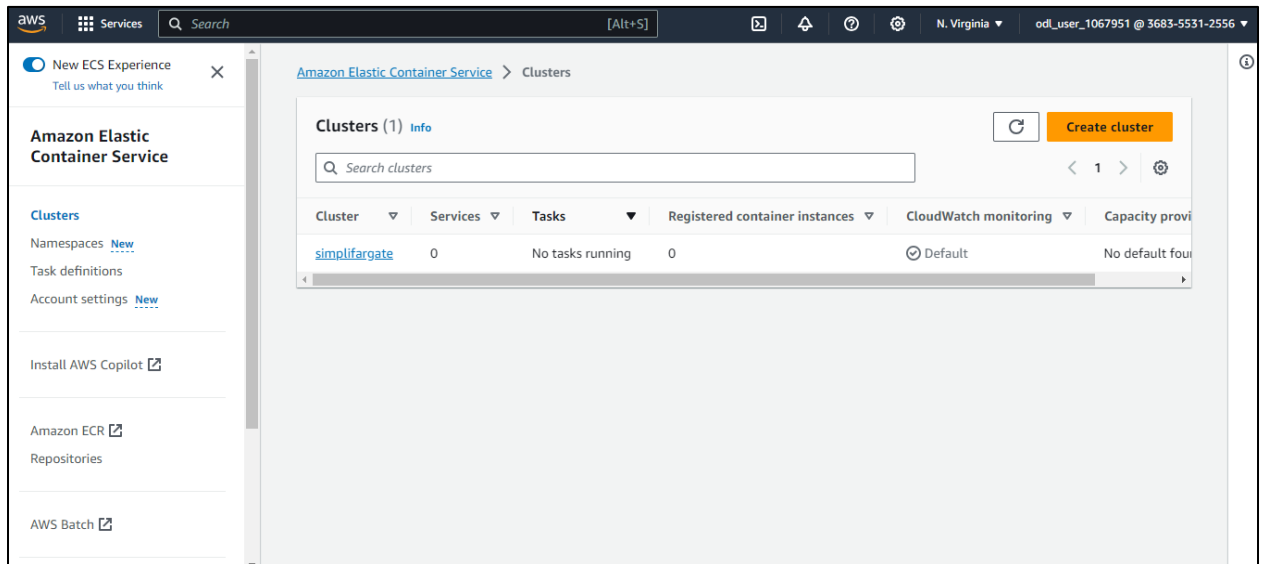
The screenshot shows the 'Task definition successfully created' notification at the top of the AWS Management Console. Below the notification, the breadcrumb trail is 'Amazon Elastic Container Service > Task definitions > nginxdef > Revision 1 > Containers'. The task definition name 'nginxdef:1' is displayed, along with 'Deploy', 'Actions', and 'Create new revision' buttons. The 'Overview' tab is selected, showing a table with the following details:

| | | | |
|---|---|---|----------------------------|
| ARN arn:aws:ecs:us-east-1:004774-726888:task-definition/nginxdef:1 | Status ACTIVE | Time created 2023-09-13T07:05:32.830Z | App environment FARGATE |
| Task role - | Task execution role ecsTaskExecutionRole | Operating system/Architecture Linux/X86_64 | Network mode awsvpc |

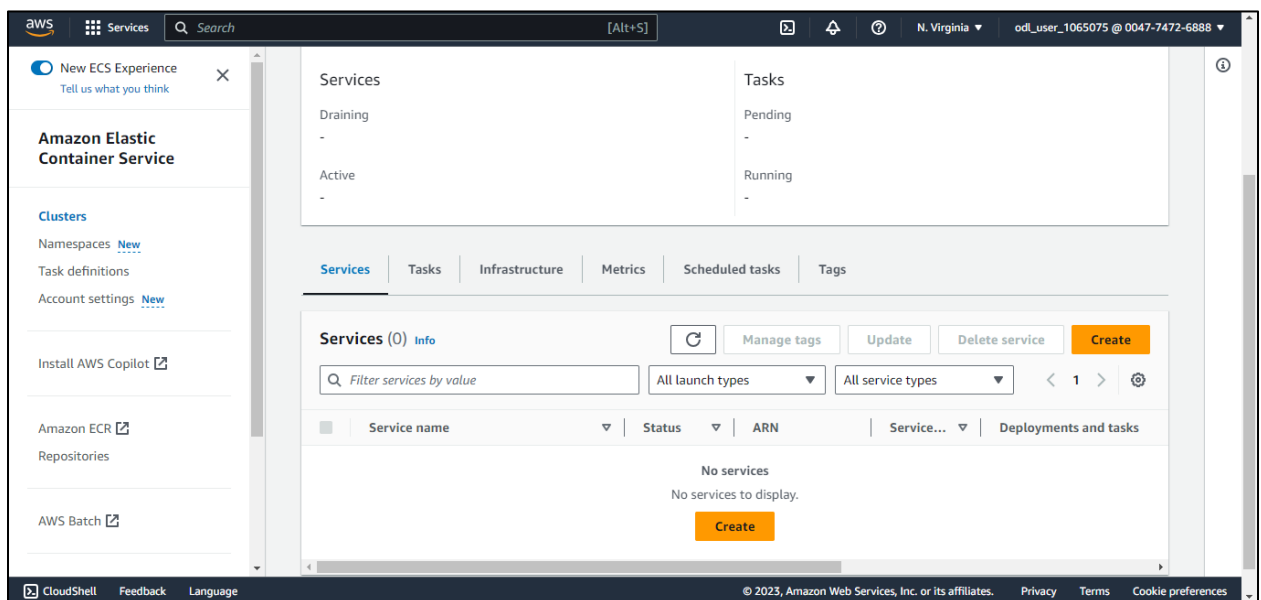
Below the table, there are tabs for 'Containers', 'JSON', 'Storage', and 'Tags'. The 'Containers' tab is currently selected.

Step 3: Run Fargate Cluster

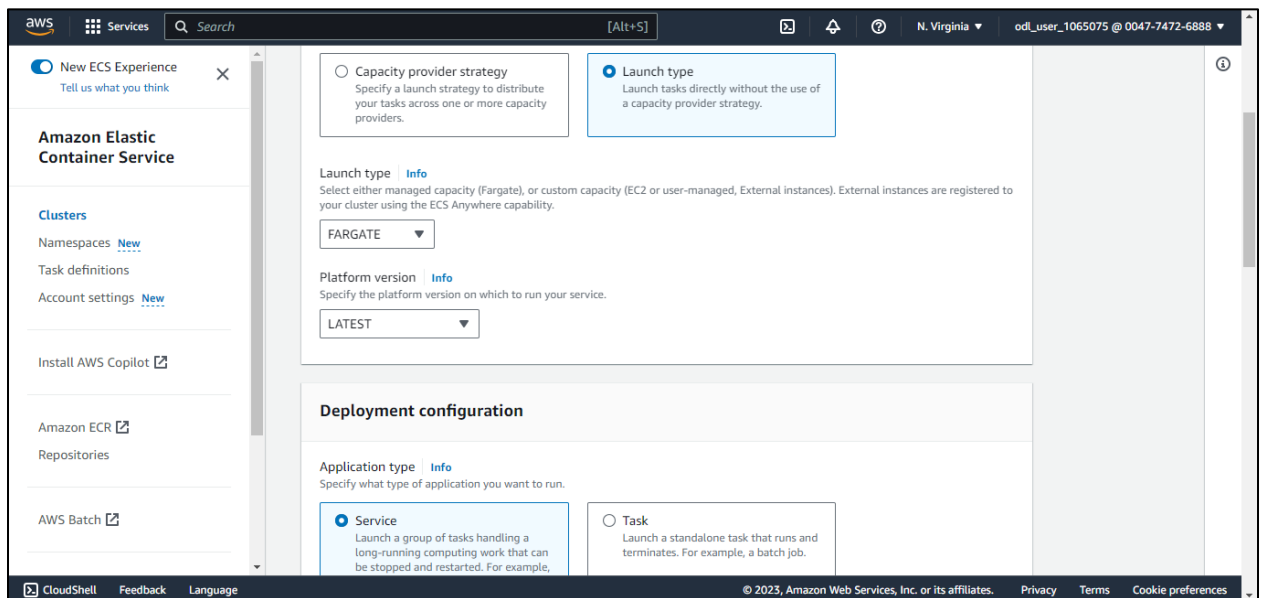
3.1 Return to the ECS home page and open the newly created cluster from **Clusters**



3.2 Click on **Create** under **Services**



3.3 Choose **FARGATE** under **Launch type**. Choose the **Platform version** as **LATEST** and the **Application type** as **Service**.



The screenshot shows the AWS Management Console interface for configuring an Amazon ECS service. The left sidebar displays the 'Amazon Elastic Container Service' navigation menu. The main content area is titled 'Launch type' and 'Platform version'. The 'Launch type' is set to 'FARGATE' and the 'Platform version' is set to 'LATEST'. Below this, the 'Application type' is set to 'Service'.

Launch type [Info](#)
Select either managed capacity (Fargate), or custom capacity (EC2 or user-managed, External instances). External instances are registered to your cluster using the ECS Anywhere capability.

Platform version [Info](#)
Specify the platform version on which to run your service.

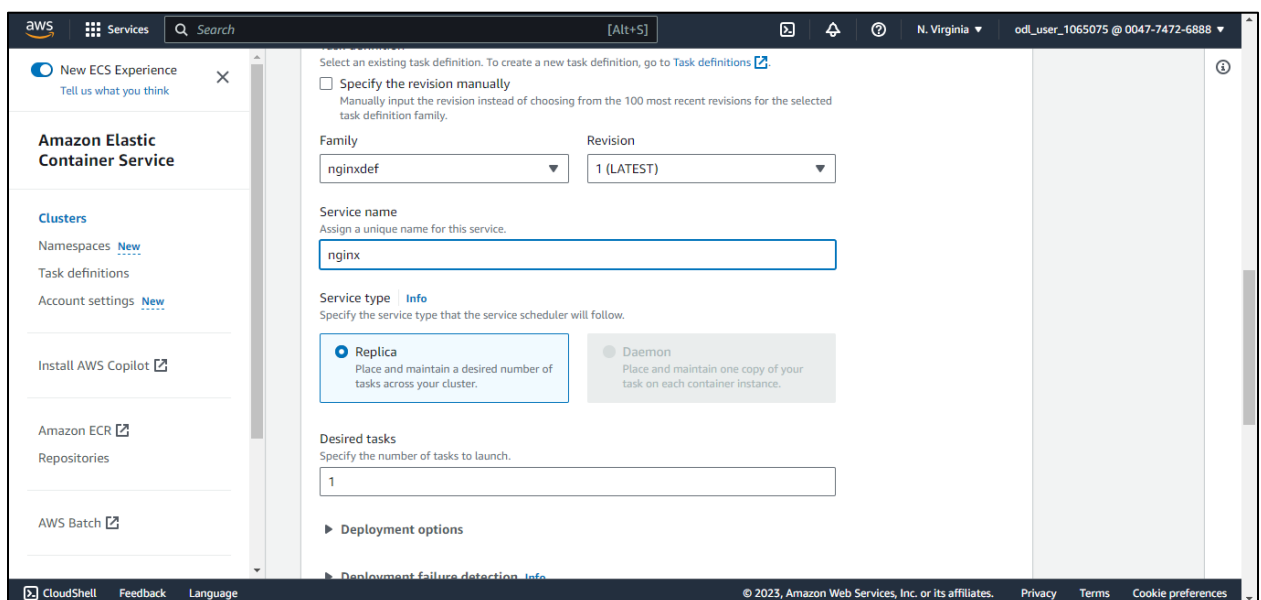
Deployment configuration

Application type [Info](#)
Specify what type of application you want to run.

Service
Launch a group of tasks handling a long-running computing work that can be stopped and restarted. For example,

Task
Launch a standalone task that runs and terminates. For example, a batch job.

3.4 Choose family = **nginxdef** (created earlier), revision = **1 (LATEST)**, service name = **nginx**, and Replica = **1**



The screenshot shows the AWS Management Console interface for configuring an Amazon ECS service. The left sidebar displays the 'Amazon Elastic Container Service' navigation menu. The main content area is titled 'Family', 'Revision', 'Service name', and 'Service type'. The 'Family' is set to 'nginxdef', the 'Revision' is set to '1 (LATEST)', the 'Service name' is 'nginx', and the 'Service type' is set to 'Replica'.

Select an existing task definition. To create a new task definition, go to [Task definitions](#).

☐ Specify the revision manually
Manually input the revision instead of choosing from the 100 most recent revisions for the selected task definition family.

Family
nginxdef

Revision
1 (LATEST)

Service name
Assign a unique name for this service.
nginx

Service type [Info](#)
Specify the service type that the service scheduler will follow.

Replica
Place and maintain a desired number of tasks across your cluster.

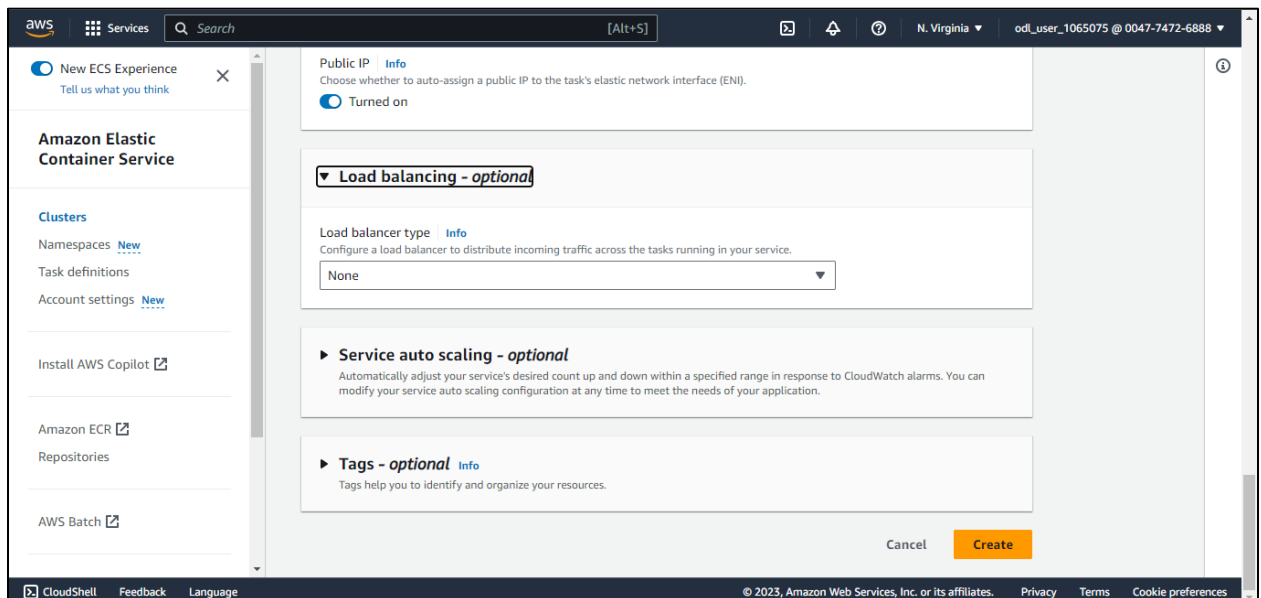
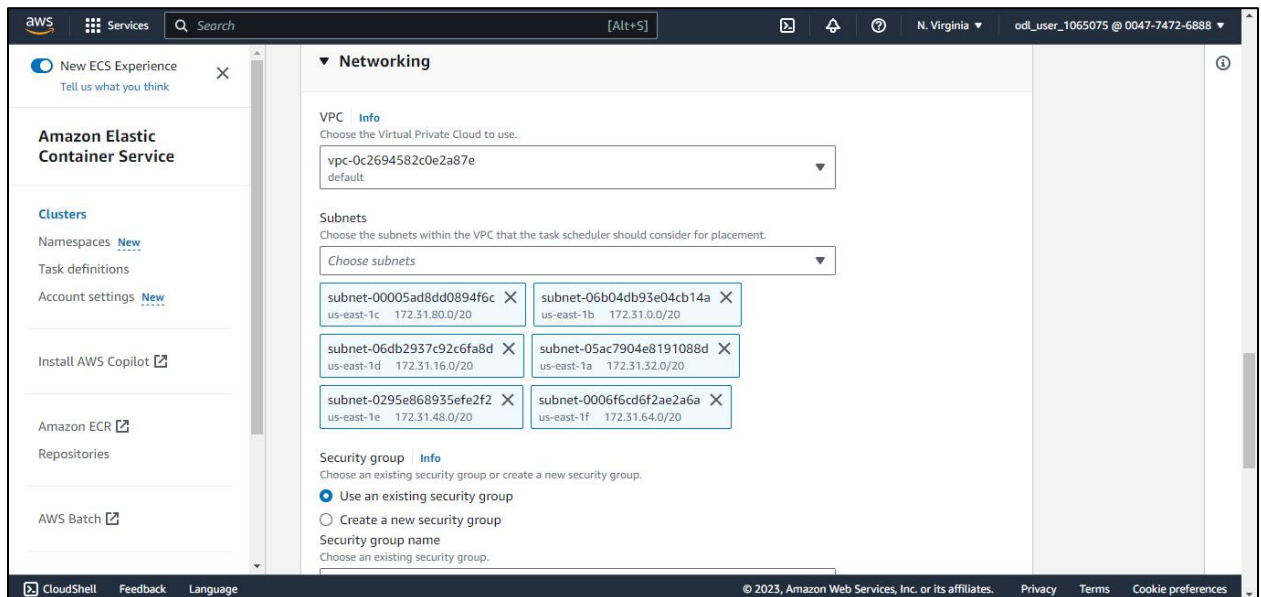
Daemon
Place and maintain one copy of your task on each container instance.

Desired tasks
Specify the number of tasks to launch.
1

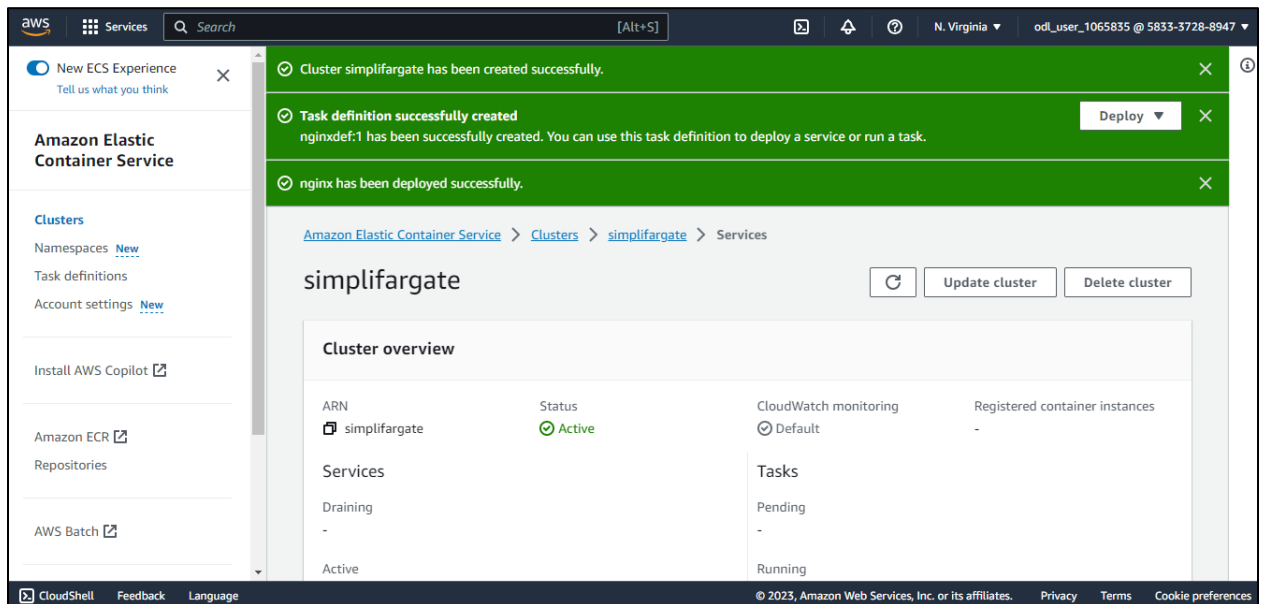
Deployment options

Deployment Failure detection [Info](#)

3.5 In the **Networking**, leave default **VPC** and **Load balancing** as **none**. Now, click **Create**.

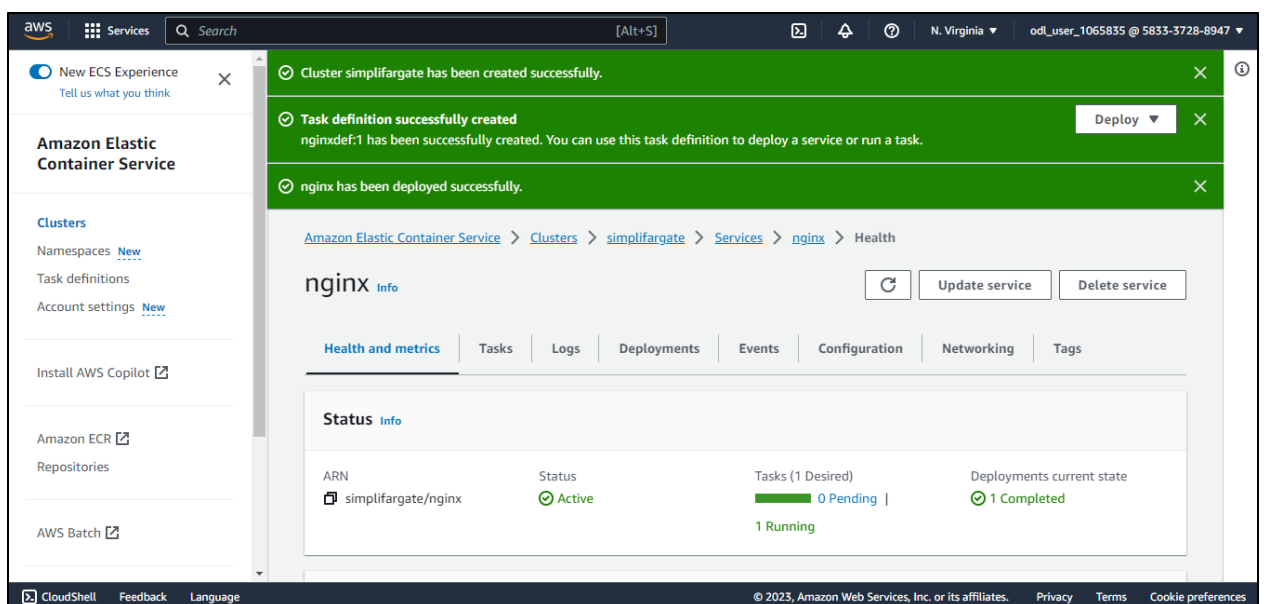


3.6 Wait until service creation is completed and 1/1 of tasks are active as shown below:

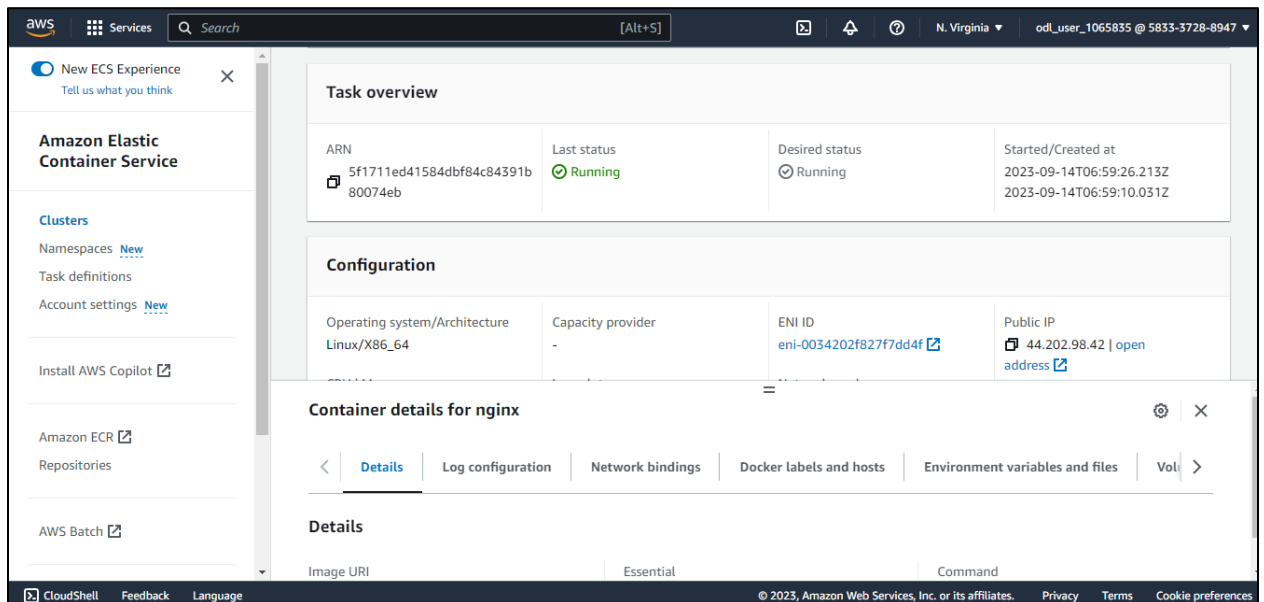


3.7 Once the service is running successfully, click **nginx** and find service details as shown below:

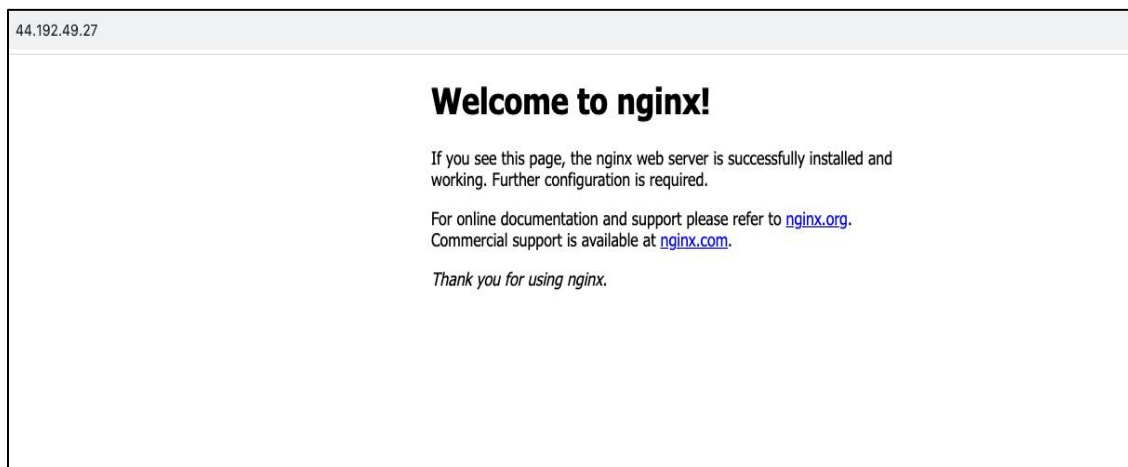
Now, click **Tasks**.



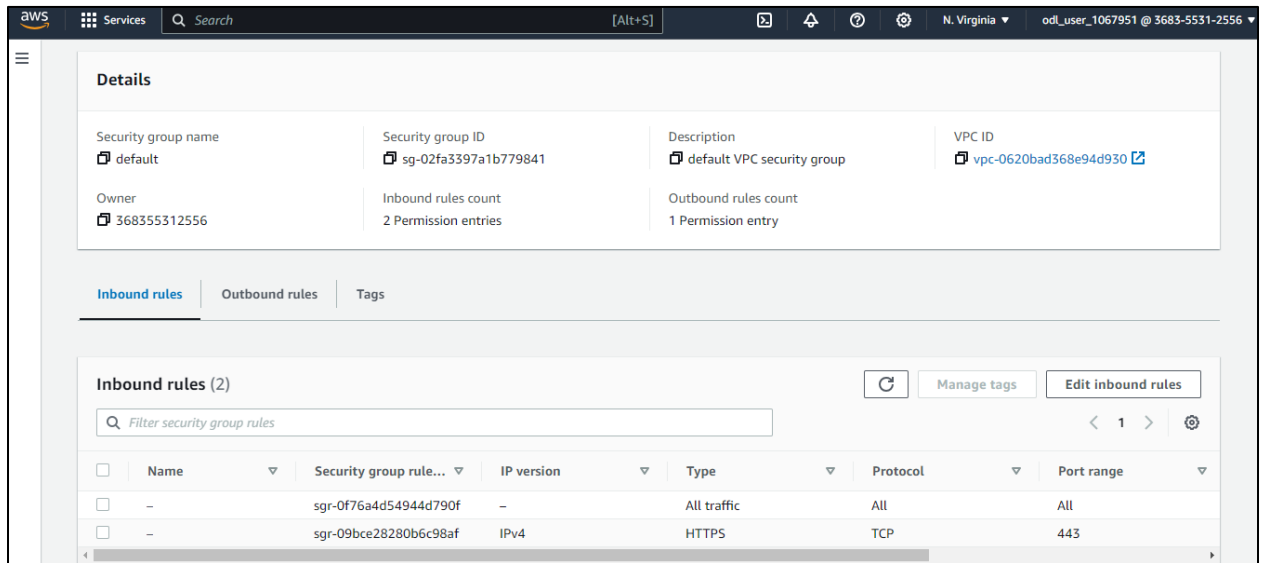
3.8 Copy the public IP address and open it in a new browser



3.9 Open the URL in the new browser to see the nginx page loading as below:



Note: In case the web page is not loading, go to **Task > Networking > Open security group**, and select the security group being used to ensure the port 80 inbound rule is allowed to be accessed from anywhere as shown below:



Hence, you have successfully executed the Fargate Cluster.